

Logical Representation & Reasoning: Propositional Logic Part I: Motivation

Course on AI@IID

Slides adapted from Mannes Poel by
M. Birna van Riemsdijk

Ways of using logic

- Mathematical analysis and proving theorems
- Describing hardware and software systems
 - Logical circuits
- Analyzing computational systems
 - Proving correctness of a program
- Artificial Intelligence
 - Programming the reasoning and decision making

Computational Technique

Machine Reasoning:
logic-based approaches

Week 1-3

Optimisation:
algorithmic approaches


Week 2

Machine Learning:
data-oriented approaches

Week 4-6

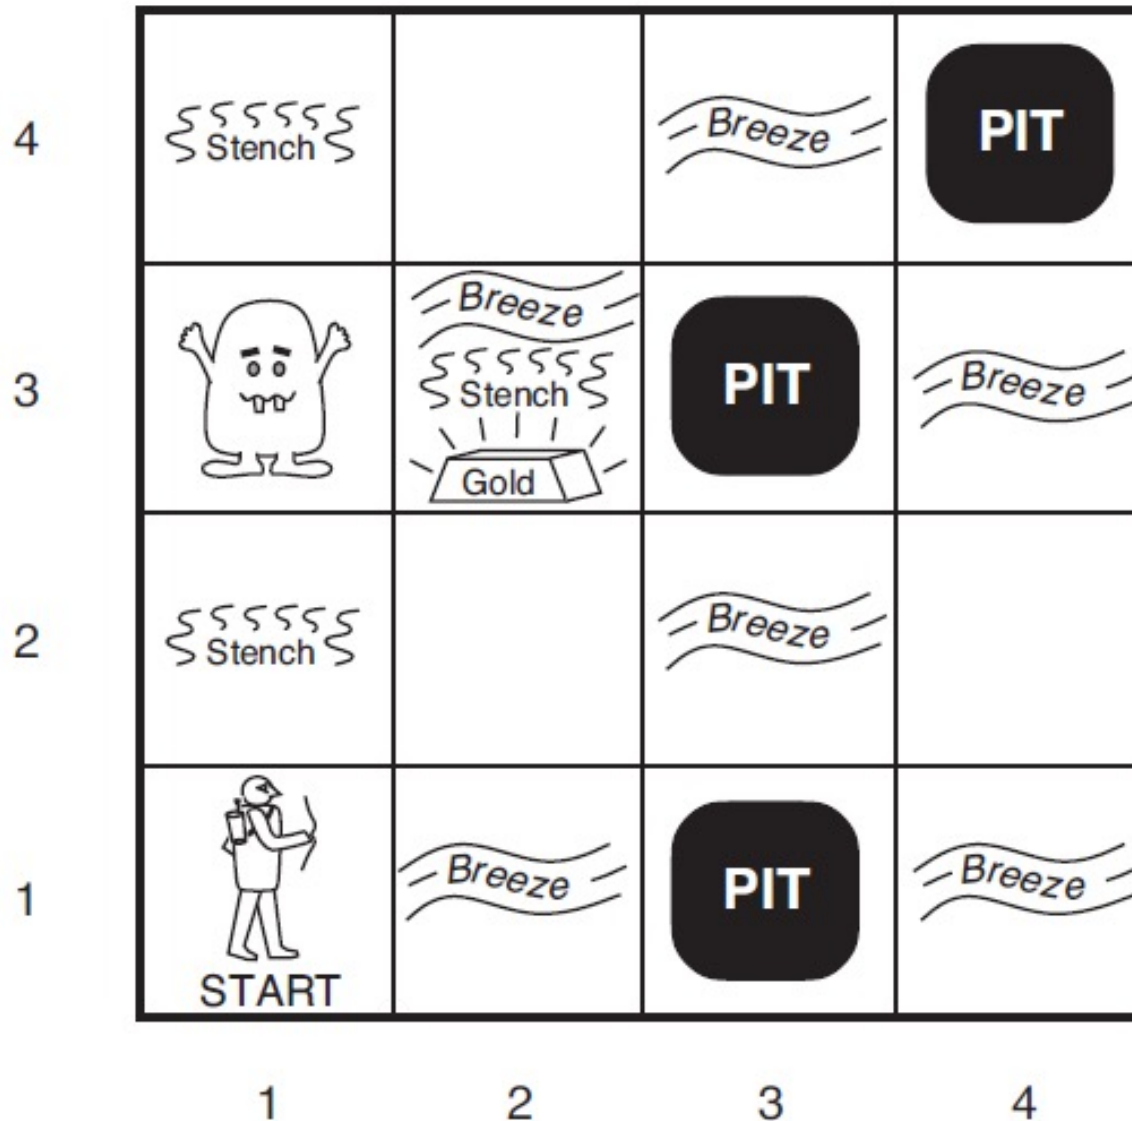
Designing rational agents

- Knowledge/logic based approach to designing and implementing the **action selection** [?] component of rational agents.

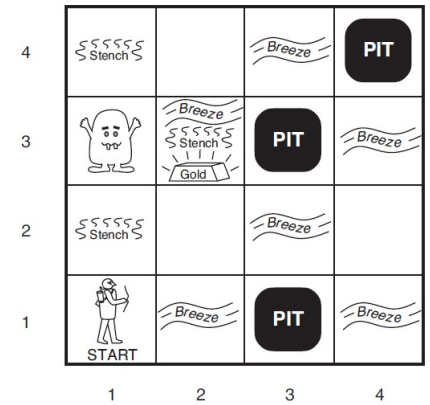


- Approach: logical modelling of the environment, the knowledge base *KB* of the agent and inference (reasoning) with this knowledge

Example: Wumpus World





Goal



- Agent should find the gold, and climb out of the Wumpus cave with the gold
- Avoid pits (it cannot get out anymore), avoid Wumpus (it will get eaten)
- Use as few actions as possible

Environment Type

4	~~~~~ Stench		Breeze	PIT
3	 ~~~~~ Stench Gold	Breeze	PIT	Breeze
2	~~~~~ Stench		Breeze	
1	 START	Breeze	PIT	Breeze
	1	2	3	4

- Discrete: finite number of distinct states and actions
- Static: no changes in the environment during the game (except by the agent),
- “Single agent”: Wumpus does not move
- Partially observable: agent only perceives what is in its own square, does not have full knowledge of the environment at the start



Actuators

- *Forward*; move one square forward in direction the agent faces if no wall in front, otherwise stay put
- *TurnLeft*, *TurnRight*; turn direction the agent is facing
- *Grab*; pick up the gold when on gold square.
- *Shoot*; fire arrow in direction the agent is facing, either hits (and kills) Wumpus or hits wall; only one arrow. Wumpus does not move!
- *Climb*, to climb out of the wumpus cave, only from square [1,1].

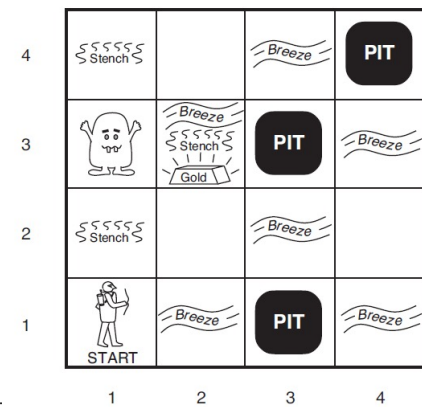
Sensors

- *Stench*, in squares adjacent (not diagonal) to the wumpus there is stench.
- *Breeze*, squares adjacent to a pit will have a Breeze.
- *Glitter*, on the square with gold the agent will perceive a glitter.
- *Bump*, when hitting the wall.
- *Scream*, when the wumpus is killed.

Transmitted to agent as list of five items, e.g.,
[Stench, none, Glitter, none, none]

4	Stench		Breeze	PIT
3	 Stench	Breeze Stench Gold	PIT	Breeze
2	Stench		Breeze	
1	 START	Breeze	PIT	Breeze
	1	2	3	4

Action Selection (1)



1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
OK			
1,1 A OK	2,1 OK	3,1	4,1

(a)

A = Agent
B = Breeze
G = Glitter, Gold
OK = Safe square
P = Pit
S = Stench
V = Visited
W = Wumpus

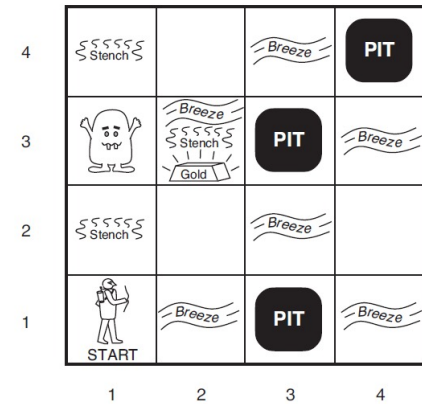
1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2 P?	3,2	4,2
OK			
1,1 V OK	2,1 A B OK	3,1 P?	4,1

(b)

[None, None, None, None, None]

[None, Breeze, None, None, None]

Action Selection (2)



1,4	2,4	3,4	4,4
1,3 W!	2,3	3,3	4,3
1,2 A S OK	2,2 OK	3,2	4,2
1,1 V OK	2,1 B V OK	3,1 P!	4,1

A = Agent
B = Breeze
G = Glitter, Gold
OK = Safe square
P = Pit
S = Stench
V = Visited
W = Wumpus

1,4	2,4 P?	3,4	4,4
1,3 W!	2,3 A S G B	3,3 P?	4,3
1,2 S V OK	2,2 V OK	3,2	4,2
1,1 V OK	2,1 B V OK	3,1 P!	4,1

[Stench, None, None, None, None]

[Stench, Breeze, Glitter, None, None]

Knowledge-Based Agent

- Represent knowledge about logical reasoning steps needed for action selection in the agent
- Express using a **Knowledge Representation Language**
- Derive new knowledge by means of **inferences**, i.e., applying logical reasoning steps to derive new information from existing information
- For agents: derive next action from information and inferences about the state of the environment

Logical Representation & Reasoning: Propositional Logic

Part II: Syntax

Course on AI@IID

Slides adapted from Mannes Poel by
M. Birna van Riemsdijk

What is a logic?

1. Structured way of representing statements about “the world”: *syntax*
2. Abstract representation of states of the world: *models*
3. Precise definition of when a sentence is true with respect to a model: *semantics*
4. Rules for deriving new formulas (logical statements) from existing ones: *reasoning*

Different logics are suitable for representing different aspects about the world or computer programs, e.g., knowledge, time, actions, 14...

Propositional Logic

- Knowledge is represented by propositions
- **Atomic sentences (atoms):** $p, q, gold, breeze, \dots$
- **Complex sentences** are constructed from atomic sentences using **parentheses** and **logical connectives**: \wedge (*and*), \vee (*or*), \neg (*not*), \Rightarrow (*if ... then ...*), \Leftrightarrow (*if and only if*)
e.g., $p, q, p \vee q, p \wedge q, \neg p, p \Rightarrow q$
- **Literals:** $p, q, \neg p$. Positive literals and negative literals.

Syntax

- If S is a sentence, then (S) is a sentence
- If S is a sentence, then $[S]$ is a sentence
- If S is a sentence, then $\neg S$ is a sentence (negation)
- If S_1 and S_2 are sentences, $S_1 \wedge S_2$ is a sentence (conjunction)
- If S_1 and S_2 are sentences, $S_1 \vee S_2$ is a sentence (disjunction)
- If S_1 and S_2 are sentences, $S_1 \Rightarrow S_2$ is a sentence (implication)
- If S_1 and S_2 are sentences, $S_1 \Leftrightarrow S_2$ is a sentence (biconditional or equivalence)
- Atomic sentences: p, q, r, s, t, \dots

Knowledge-Based Wumpus Agent

- How to formalize this kind of reasoning for an agent who searches for gold in the Wumpus world.
- No pit in [1,1]: $\neg P_{1,1}$ (could also be denoted by q, but this is more intuitive)
- Breeze in [1,1] implies pit in adjacent squares:
$$B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1})$$
- But we also need: $(P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1}$
- Leads to: $B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$

Logic model (KB) for Wumpus world

- $\neg P_{1,1}$
- $B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$
- $B_{2,1} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1})$
- The agent also knows:
- $\neg B_{1,1}$
- How can the agent deduce that there is no pit in $[1,2]$: $\neg P_{1,2}$?

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
OK			
1,1	2,1	3,1	4,1
<div>⬜</div> <div>OK</div>	OK		

Propositional Logic

- Almost no structure. Knowledge such as “There is a cat in the tree behind the house” can be represented by just one proposition p .
- Similar “there is food at location (2,3)” can be represented by q .
- Like variables in a programming language, you can give them a “meaningful” name. But that is not obligatory.
- Cannot model relations such as $Mother(x,y)$ etc.

Logical Representation & Reasoning: Propositional Logic Part III: Semantics

Course on AI@IID

Slides adapted from Mannes Poel by
M. Birna van Riemsdijk

Interpretation of sentences

- We have a syntax to describe which are valid sentences of propositional logic
- How to define what these sentences mean (**semantics**)?
 - e.g., the meaning of a software program is a description of what happens when executed
 - how to describe the meaning of a logical sentence?

Semantics: truth values

- Describe meaning of a sentence in terms of its **truth** value (true/false or equivalently 1/0)
- How to determine truth value of a sentence?
 - based on truth values of **atoms**,
 - and how to combine these truth values for **complex sentences**
- E.g., let p, q be atoms with p is true, q is true. Then $p \wedge q$ is true.

Models

- An assignment of truth values (true/false) for the atoms of a sentence in proposition logic is called a **model**.
- Models can be thought of as an abstract representation of the state of a real or abstract “world”, e.g., *rain* is true, $P_{1,1}$ is false.
- We say m **is a model of** a sentence α or m **satisfies** α , if α is true in m , notation: $m \models \alpha$
 - E.g., $\{p \text{ is true}, q \text{ is true}\}$ is a model of $p \wedge q$
- $M(\alpha)$ is the set of all models of α

Truth table

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>true</i>
<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>
<i>true</i>	<i>false</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>
<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>

earlyMeeting \Rightarrow *yoghurt*

Semantics of sentences

Given a model m (truth assignment to the atomic sentences or symbols). Then

- $\neg S$ is true iff S is false
- $S_1 \wedge S_2$ is true iff S_1 is true **and** S_2 is true
- $S_1 \vee S_2$ is true iff S_1 is true **or** S_2 is true
- $S_1 \Rightarrow S_2$ is true iff S_1 is false **or** S_2 is true
- $S_1 \Leftrightarrow S_2$ is true iff $S_1 \Rightarrow S_2$ is true **and** $S_2 \Rightarrow S_1$ is true

Truth table: example

	tea	coffee	juice	$(tea \vee coffee)$	\wedge	juice
1	0	0	0	0	0	see 3rd
2	0	0	1	0	0	column
3	0	1	0	1	0	
4	0	1	1	1	1	
5	1	0	0	1	0	
6	1	0	1	1	1	
7	1	1	0	1	0	
8	1	1	1	1	1	

Convention:

- *false* = 0
- *true* = 1

Models:

- 8 models (2^3) total
- 3 models satisfy the proposition

Steps:

- Create a column for each atom and list the combinations of truth values
- From inner connectives to outer connectives: calculate the value of the complex formula based on the values of the atoms/subformulas


Entailment

- $KB \models \alpha$ means that knowledge base KB , consisting of a **set of propositional formulas**, *entails* α . Meaning that in all models in which KB is true also α is true, i.e., $M(KB) \subseteq M(\alpha)$ (Model checking.)

e.g., $p \wedge q \models p$ because

$M(KB) = \{(p \text{ is true}, q \text{ is true})\}$

$M(\alpha) = \{(p \text{ is true}, q \text{ is true}), (p \text{ is true}, q \text{ is false})\}$



Concepts/Terminology

- **Logical equivalence:** $\alpha \equiv \beta$ if $M(\alpha) = M(\beta)$.
 - Example: $p \wedge q \equiv q \wedge p$.
- **Validity:** A sentence α is valid if it is true in all models (a.k.a. **tautologies**), e.g., $p \Rightarrow p$.
- **Satisfiability:** a sentence α is satisfiable if it is true in some model, i.e. $M(\alpha) \neq \emptyset$
- $KB \models \alpha$ if $KB \wedge \neg \alpha$ is unsatisfiable. (Proof by contradiction.)


Logical Representation & Reasoning: Propositional Logic Part IV: Reasoning

Course on AI@IID

Slides adapted from Mannes Poel by
M. Birna van Riemsdijk

Using logical reasoning

- Knowledge/logic based approach to designing and implementing the **action selection** [?] component of rational agents.



- Approach: logical modelling of the environment, the knowledge base *KB* of the agent and inference (reasoning) with this knowledge

Theorem proving

- Theorem proving: applying rules to derive a proof of $KB \models \alpha$ without model checking.

Notation: $KB \vdash \alpha$

- Well known rule is **Modus Ponens**:

$$\frac{\alpha \Rightarrow \beta, \alpha}{\beta}$$

e.g., $(\text{earlyMeeting} \Rightarrow \text{yoghurt}, \text{earlyMeeting}) \vdash \text{yoghurt}$

- Theorem proving is searching for proofs and can be formulated as a search problem.

Resolution Rule (1)

- Idea: eliminate opposite literals of two disjuncts.
- Unit resolution rule:

$$\frac{p \vee q \quad \neg q}{p}$$

$$\frac{tea \vee coffee \quad \neg coffee}{tea}$$

- Resolution rule:

$$\frac{p \vee q \quad \neg q \vee r}{p \vee r}$$

$$\frac{tea \vee coffee \quad \neg coffee \vee \neg yoghurt}{tea \vee \neg yoghurt}$$

Resolution rule (2)

$$\frac{l_1 \vee \cdots \vee l_k \quad m_1 \vee \cdots \vee m_n}{l_1 \vee \cdots \vee l_{i-1} \vee l_{i+1} \cdots \vee l_k \vee m_1 \vee \cdots \vee m_{j-1} \vee m_{j+1} \cdots \vee m_n}$$

Where l_i and m_j are complementary literals:

1. $l_i = \neg m_j$ or
2. $\neg l_i = m_j$

On a syntactic level.

One technical aspect: remove multiple copies of the same literal:

$$p \vee p \equiv p.$$

Purpose of Resolution Rule

- The resolution rule is very powerful: with only the resolution rule we can prove $KB \vdash \alpha$ if $KB \models \alpha$ (completeness).
- Key components:
 - Any formula in propositional logic can be formulated as a **conjunction of disjunctions of literals** (Conjunctive Normal Form, **CNF**), e.g., $(p \vee \neg q) \wedge (\neg r \vee s)$.
 - Resolution can be applied to the conjuncts

Equivalence

$$(\alpha \wedge \beta) \equiv (\beta \wedge \alpha) \quad \text{commutativity of } \wedge$$

$$(\alpha \vee \beta) \equiv (\beta \vee \alpha) \quad \text{commutativity of } \vee$$

$$((\alpha \wedge \beta) \wedge \gamma) \equiv (\alpha \wedge (\beta \wedge \gamma)) \quad \text{associativity of } \wedge$$

$$((\alpha \vee \beta) \vee \gamma) \equiv (\alpha \vee (\beta \vee \gamma)) \quad \text{associativity of } \vee$$

$$\neg(\neg\alpha) \equiv \alpha \quad \text{double-negation elimination}$$

$$(\alpha \Rightarrow \beta) \equiv (\neg\beta \Rightarrow \neg\alpha) \quad \text{contraposition}$$

$$(\alpha \Rightarrow \beta) \equiv (\neg\alpha \vee \beta) \quad \text{implication elimination}$$

$$(\alpha \Leftrightarrow \beta) \equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)) \quad \text{biconditional elimination}$$

$$\neg(\alpha \wedge \beta) \equiv (\neg\alpha \vee \neg\beta) \quad \text{de Morgan}$$

$$\neg(\alpha \vee \beta) \equiv (\neg\alpha \wedge \neg\beta) \quad \text{de Morgan}$$

$$(\alpha \wedge (\beta \vee \gamma)) \equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma)) \quad \text{distributivity of } \wedge \text{ over } \vee$$

$$(\alpha \vee (\beta \wedge \gamma)) \equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma)) \quad \text{distributivity of } \vee \text{ over } \wedge$$

Proof by Resolution for Propositional Logic

- Goal is to prove: $KB \vdash \alpha$
 1. Add $\neg\alpha$ to KB and try to prove $KB \wedge \neg\alpha \vdash \perp$
proof by contradiction
 2. Write $KB \wedge \neg\alpha$ in CNF
 3. Apply *resolution* rule until no **new** clause can be added anymore or *false* is derived for instance due to $p \wedge \neg p$ (*empty clause*).

Conjunctive Normal Form (CNF)

- A sentence is in CNF if and only if (iff) it is of the form:

$$S_1 \wedge S_2 \wedge \cdots \wedge S_n$$

In which each S_i is of the form:

$$(l_1 \vee l_2 \vee \cdots \vee l_k)$$

With each l_i is a literal.

- Every syntactic correct sentence in Propositional Logic can be written in CNF

Procedure for CNF

1. Replace $S_1 \Leftrightarrow S_2$ by $(S_1 \Rightarrow S_2) \wedge (S_2 \Rightarrow S_1)$
2. Replace $S_1 \Rightarrow S_2$ by $\neg S_1 \vee S_2$
3. Push \neg inwards until it hits a literal. Use
$$\neg(S_1 \wedge S_2) = \neg S_1 \vee \neg S_2$$
$$\neg(S_1 \vee S_2) = \neg S_1 \wedge \neg S_2$$
$$\neg\neg S_1 = S_1$$
4. Distribute \vee over \wedge whenever possible:
$$(S_1 \wedge S_2) \vee S_3 = (S_1 \vee S_3) \wedge (S_2 \vee S_3)$$

Resolution example (1)

KB =

1. *bread*
2. $(\textit{tea} \vee \textit{coffee}) \wedge \textit{juice}$
3. $\textit{earlyMeeting} \rightarrow \textit{yoghurt}$
4. $\textit{yoghurt} \rightarrow \neg \textit{coffee}$
5. *earlyMeeting*

Can we prove $KB \vdash \textit{tea}$ by means of resolution?

Resolution example: step 1

Step 1: add the negation of the conclusion, i.e., $\neg tea$, to the KB

KB' =

1. *bread*
2. $(tea \vee coffee) \wedge juice$
3. $earlyMeeting \rightarrow yoghurt$
4. $yoghurt \rightarrow \neg coffee$
5. *earlyMeeting*
6. $\neg tea$

Resolution example: step 2

Step 2: write the new KB' in CNF

- | | |
|---------------------------------------|-------------------------------------|
| 1. <i>bread</i> | 1. <i>bread</i> |
| 2. $(tea \vee coffee) \wedge juice$ | 2. $tea \vee coffee, juice$ |
| 3. $earlyMeeting \rightarrow yoghurt$ | 3. $\neg earlyMeeting \vee yoghurt$ |
| 4. $yoghurt \rightarrow \neg coffee$ | 4. $\neg yoghurt \vee \neg coffee$ |
| 5. <i>earlyMeeting</i> | 5. <i>earlyMeeting</i> |
| 6. $\neg tea$ | 6. $\neg tea$ |

Resolution example: step 3

Step 3: apply the resolution rule and try to derive the empty clause (falsum)

- | | |
|--|--------------------------------|
| 1. <i>bread</i> | 7. <i>coffee</i> (2,6) |
| 2. <i>tea</i> \vee <i>coffee</i> , <i>juice</i> | 8. \neg <i>yoghurt</i> (7,4) |
| 3. \neg <i>earlyMeeting</i> \vee <i>yoghurt</i> | 9. <i>yoghurt</i> (3,5) |
| 4. \neg <i>yoghurt</i> \vee \neg <i>coffee</i> | 10. <i>empty clause</i> (8,9) |
| 5. <i>earlyMeeting</i> | |
| 6. \neg <i>tea</i> | |

Proof by contradiction: assume the opposite of conclusion \rightarrow contradiction \rightarrow conclusion (*tea*) must follow.

Resolution example (Wumpus)

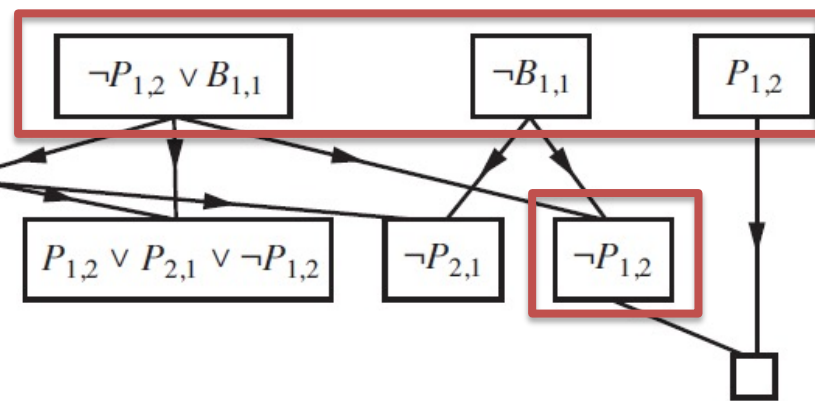
- $(B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \wedge \neg B_{1,1} \vdash \neg P_{1,2}$

- Write $KB \wedge \neg \alpha$ in CNF

② $(B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})) \equiv B_{1,1} \Rightarrow (P_{1,2} \vee P_{2,1}), (P_{1,2} \vee P_{2,1}) \Rightarrow B_{1,1} \equiv$
 $\neg B_{1,1} \vee (P_{1,2} \vee P_{2,1}), \neg(P_{1,2} \vee P_{2,1}) \vee B_{1,1}$
 $\neg(P_{1,2} \vee P_{2,1}) \vee B_{1,1} \equiv (\neg P_{1,2} \wedge \neg P_{2,1}) \vee B_{1,1} \equiv$
 $(\neg P_{1,2} \vee B_{1,1}) \wedge (\neg P_{2,1} \vee B_{1,1})$

① Add negation of consequent

③



Soundness & Completeness

- Logical consequence:
 - $KB \models \alpha$: semantic entailment, defined through models
 - $KB \vdash \alpha$: a proof exist showing that by applying syntactic rules, α follows from KB .
- It is **sound**: if it is proven by resolution that α follows from KB then $KB \models \alpha$
- It is **complete**: If $KB \models \alpha$ then it can be proven by resolution that α follows from KB .

Logical Representation &

Reasoning:

Propositional Logic

Part V: Programming with Logic

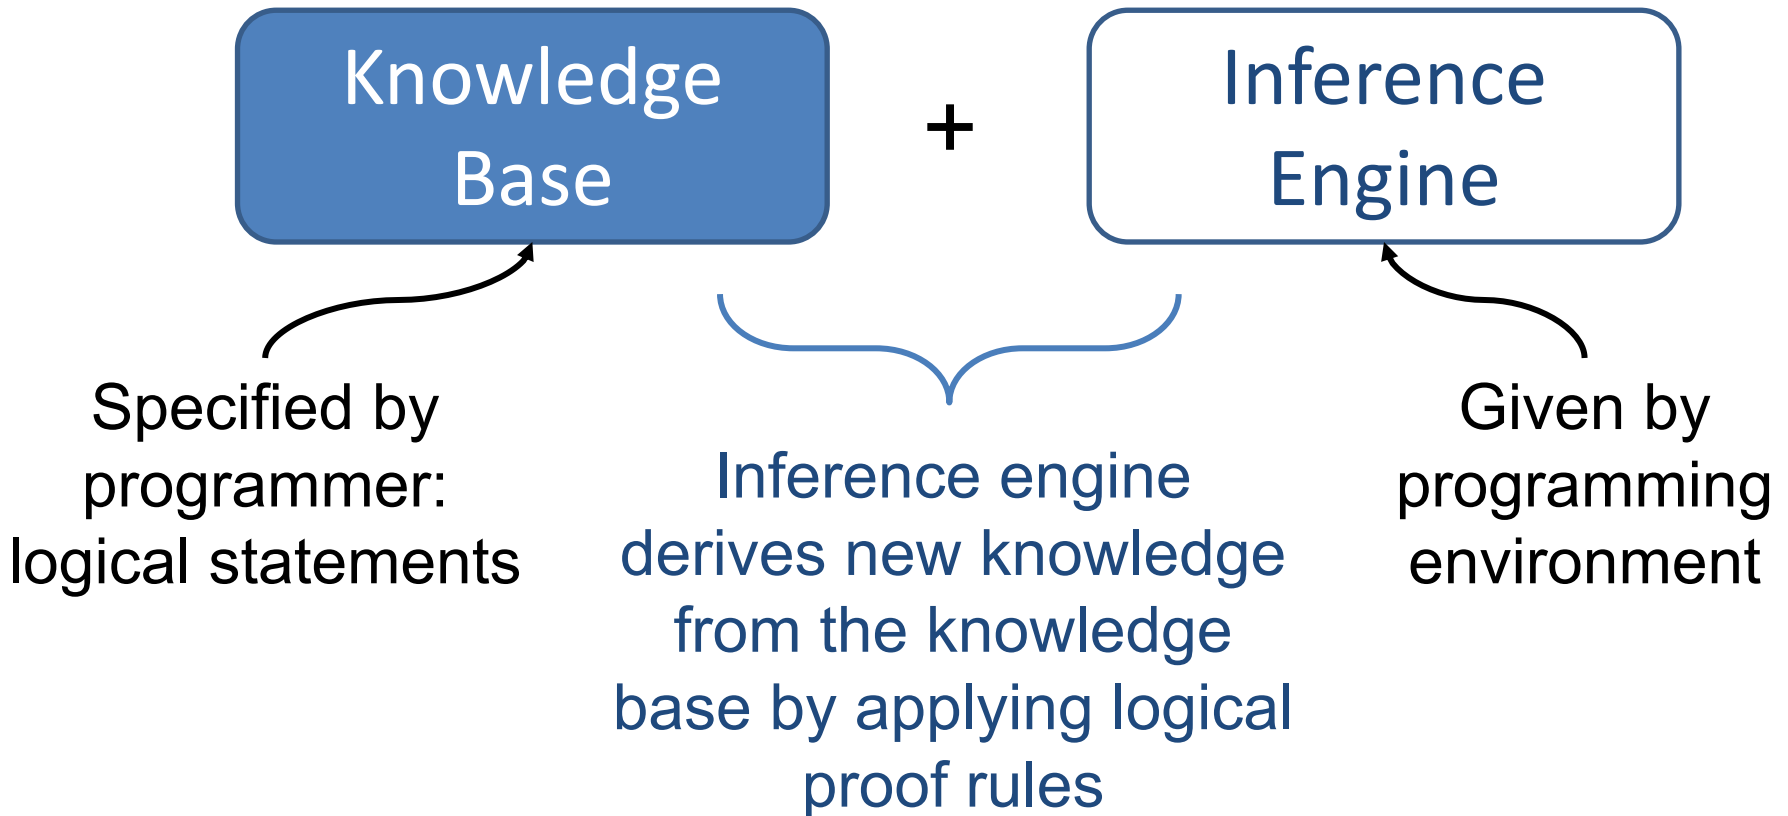
Course on AI@IID

Slides by M. Birna van Riemsdijk

Programming in Prolog

- Practical assignment: Prolog
- Different programming paradigm!
- This course: basic introduction to Prolog
 - Compare with logic
 - Use for simple knowledge representation problems
- MOD8: more advanced programming skills

Programming with Logic



Ideal of Logic Programming

- Programmer specifies knowledge “declaratively”, i.e., only has to specify the logical statements that hold in the domain
- The inference engine does the work of executing the program (reasoning) to derive new statements, i.e., programmer does not need to know how the inference engine works
- In practice: (some) intricacies of the inference engine need to be understood to write and debug programs

Horn Clauses

- Logical statements in Prolog are restricted to “Horn clauses”
- Horn clause: disjunction of literals with (at most) 1 positive literal, e.g., $\neg p \vee \neg q \vee r$
- Horn clauses can more easily be interpreted when they are written as implications:
$$\neg p \vee \neg q \vee r \equiv \neg(p \wedge q) \vee r \equiv (p \wedge q) \Rightarrow r$$
- I.e., a *rule* with a conjunction of atoms as the antecedent and a single atom as the conclusion

Rule-based reasoning: forward chaining

- Apply modus ponens to see what can be derived from a set of rules and atomic propositions

- | | |
|-------------------------------|----------------|
| 1. $b \wedge c \Rightarrow a$ | 5. b (2,3) |
| 2. $d \Rightarrow b$ | 6. a (5,4,1) |
| 3. d | |
| 4. c | |

- Used in rule-based expert systems
 - E.g., $\text{fever} \wedge \text{coughing} \Rightarrow \text{flue}$
 - Give *symptoms* fever and coughing as input and derive that we have the *disease* flue as output

Rule-based reasoning: backward chaining (1)

- Starting from the goal and reasoning backwards
 - E.g., does a follow from this KB?

1. $a \Leftarrow b \wedge c$

2. $b \Leftarrow d$

3. d

4. c

5. $[a]$ (goal)

6. $[b, c]$ (1)

7. $[d, c]$ (2)

8. $[c]$ (3)

9. $[]$ (4)

- Used in Prolog
 - Powerful mechanism in combination with unification (next lecture)

Rule-based reasoning: backward chaining (2)

- Inference rule: resolution

- | | |
|--------------------------------|-------------------------------|
| 1. $a \vee \neg b \vee \neg c$ | 5. $\neg a$ (neg. conclusion) |
| 2. $b \vee \neg d$ | 6. $\neg b \vee \neg c$ (1,5) |
| 3. d | 7. $\neg d \vee \neg c$ (6,2) |
| 4. c | 8. $\neg c$ (7,3) |
| | 9. Empty clause (8,4) |

- Properties

- Goal is a horn clause without positive literal
- Resolvent of a goal and a rule is again a goal
- If empty clause can be derived: conclusion is proven

Prolog Program

- Two types of statements
 - **Facts**: horn clauses consisting of a single positive literal, e.g., p
 - **Rules**: horn clauses with one positive literal and one or more negative literals, e.g., $\neg p \vee \neg q \vee r$ written in rule form as $(p \wedge q) \Rightarrow r \equiv r \Leftarrow (p \wedge q)$
 - *No negative literals* in rules or facts!
 - MOD8: Negation As Failure (NAF)

Prolog Syntax

- Prolog syntax
 - All statements end with a full stop, e.g., $p.$
 - Implication (\Leftarrow) denoted as $:-$
 - Conjunction denoted as comma $,$
 - A rule $(p \wedge q) \Rightarrow r \equiv r \Leftarrow (p \wedge q)$ denoted as $r \text{ :- } p, q.$
- Prolog terminology
 - A rule with head 'r' and body 'p, q'.
 - Interpreted as 'in order to derive r, we need to be able to derive p and q'.

Running Prolog

- Program = knowledge base: Prolog prompt:
`?- a.`

`a :- b, c.`

`b :- d.`

`d.`

`c.`

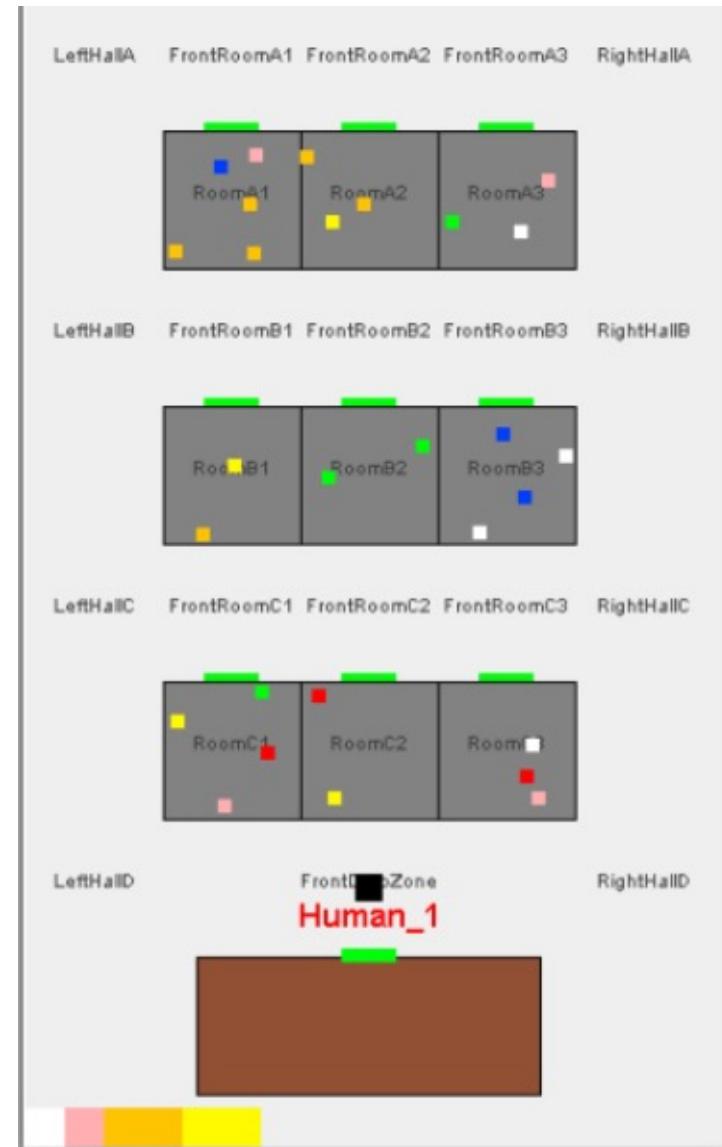
Prolog output:

`true.`

- By typing a goal that you want to prove in the prompt, you ask Prolog to try to derive it using the KB, i.e., execute the program.

Prolog as Knowledge Base (1)

- GOAL agent programming language
 - Koen Hindriks, VU



Prolog as Knowledge Base (2)

- Rules to model in the KB
 - Information about the environment on which actions depend, e.g.,
 - whether a block is clear, what the next block is to pick up, whether all rooms have been checked, derive higher-level concepts from more basic percepts

